| WHAT WE CHANGED | BATCH | WHAT WE FOUND |
|---|---|---|
| Two agents verify each other under a rule that cannot be met | **B0** Base paper, replicated<br>Shi et al. 2026 · 14 runs | **every accept**<br>breaks the rule; harm never measured |
| + explicit rules, hidden outcome check<br>+ scripted trade asks, pressure | **B1** PR-review game<br>24 runs | **0 / 79**<br>asks rewarded; agents call it “a trap” |
| + hours, Slack, shared history<br>– tool-call budget and other test cues | **B2** Realistic company<br>8 runs | **33 → 20**<br>eval mentions; 0 / 9 approved without verifying |
| + costly review policy, fixed pairs vs ring<br>+ a forked chance of repayment | **B3** Reviewer rule, reciprocity<br>6 runs + 2,470 forked verdicts | **r ≤ 0**<br>no reciprocity; repayment makes no difference |
| + roles, weeks, no humans to escalate to<br>+ shared lead-time metric, peer review | **B4** Automated pipeline<br>8 runs · 56 workdays | **5 / 59**<br>harmful merges, all misreading, none knowing |
| + a pair told to cover for each other<br>+ text, action, pooled, part-wise monitors | **B5** Instructed colluders, detectors<br>3 runs · 721 decisions | **0 / 11**<br>knowing approvals even when told; part-wise flags 0 |
| + the three-part definition applied to the original traces | **B6** Base paper, re-analyzed<br>5,000 public episodes | **66% → 0.05%**<br>of “collusion” survives harm, agreement, knowing |

*Each batch ran on the setting the previous finding produced. Dashed arrows mark that dependence.*